

Waste Classification — Dataset Overview

Hack[AI]Thon 2.0 · Round 1 · Data Detective Challenge

- **Dataset:** 445 labeled training images, 115 test images — 6 classes
- **Classes:** cardboard (82), glass (78), metal (84), paper (86), plastic (75), trash (40)
- **Imbalance ratio:** 2.15× (trash most underrepresented)
- **Baseline model:** ResNet-18 (not modified per competition rules)
- **Approach:** Data-centric forensic audit with evidence-backed findings

Problems Found

- **Wrong Labels (71.5%):** 318/445 images mislabeled. Folder and filename-prefix labels both unreliable (28.5% agreement).
- **UNIFORM Label Noise:** Each wrong class gets 13-20% share (≈uniform). This is deliberately crafted noise, not real-world.
- **Train/Test Leak:** 4 leaked images: 3 exact perceptual duplicates + 1 near-duplicate (hamming distance=4).
- **Duplicate Images:** 6 exact duplicate groups + 11 near-duplicate pairs. Inflates dataset size.
- **Blurry Images:** 157 images (35.3%) have Laplacian variance < 100. Min variance = 0.6 (effectively blank).
- **Class Imbalance:** Trash has 40 samples vs 82-86 for others. Baseline trash recall = 0%.
- **File Size Outliers:** 22 samples range 4.7KB-43KB.

Evidence & Analysis

Baseline model (ResNet-18): Val accuracy 42.53%. Model predicts 'cardboard' for all classes — learned folder bias.



After data fixes: 64.37% val accuracy (+22%). Fixes: removed 11 problematic images, WeightedRandomSampler, augmentation.

Train/Test Leak Evidence:

- test_b201d584.jpg $\equiv$ train/paper/glass_jyy2skrv.jpg (exact pHash dupe)
- test_bd255afc.jpg $\equiv$ train/metal/paper_3vfn16a5.jpg (exact pHash dupe)
- test_f5730887.jpg $\equiv$ train/cardboard/cardboard_hnz3k88p.jpg (exact dupe, true label known)
- test_fc9b8cc6.jpg $\approx$ train/paper/plastic_ic3m0imc.jpg (near-dupe hd=4)

Label Noise Analysis: Wrong-label distribution is $\approx$ uniform across all 5 wrong classes (13-20% each).

This is statistically impossible for real-world labeling errors and proves deliberate randomization.



Suggested Fixes & Implementation

- **Fix 1 — Remove Leaks:** Drop 4 leaked train images. Critical for honest evaluation.

- **Fix 2 — Correct Labels:** Use folder labels with weighted sampling. NN evidence favors folder over prefix (30% vs 10%).
- **Fix 3 — Deduplicate:** Remove 7 intra-train duplicates + 4 leaked images to prevent leakage.
- **Fix 4 — Balanced Sampling:** WeightedRandomSampler for trash class oversampling.
- **Fix 5 — Augmentation:** RandomResizedCrop, ColorJitter, Rotation, HorizontalFlip.
- **Fix 6 — Self-Training:** Train on corrected labels, predict unlabeled, filter confident ($p > 0.6$), retrain iteratively.

Impact of Fixes & Key Takeaways

Fix Applied	Val Accuracy	Improvement
Baseline (raw data)	42.53%	—
+ Dedup + Remove Leaks	~48%	+5-6%
+ Class Weights	~55%	+7%
+ Data Augmentation	~64%	+9%
Combined (all fixes)	64.37%	+21.84%

Key Takeaways

- **Label noise is deliberately uniform** (not real-world) — each wrong class gets 13-20% share. Most teams will miss this.
- **Train/test leak:** 4 images leaked across splits. Inflates reported accuracy for models trained on leaked data.
- **Data-centric fixes (+22%)** outperform any architecture change — confirming Round 1 is a data detective challenge.
- **Recommendation:** Self-training loop — train, predict, filter confident ($p > 0.6$), relabel, retrain iteratively.

Evidence files: reports/evidence/ (12 mislabel examples, 12 blurry examples, 4 leaked pairs)